DiffCP: Full Technical Analysis

From Proposal to Working Solver — Complete implementation, benchmark results, and roadmap for CS257

- **From:** Rick Underwood
- **Date:** March 16, 2026
- **For:** Vijay Daita · Hayley Antczak · Miro Swisher
- **Course:** CS257

SECTION 1

What We Built

Starting from the one-page proposal, we built a complete differentiable constraint programming solver — a working research prototype demonstrating that gradient-based methods can handle combinatorial satisfaction problems at practical scale.

PACKAGE

diffcp/ Python package — 7 modules, 5,000+ lines across solver, rounding, constraint, and model layers.

API

CP-SAT compatible interface: drop-in replacement for the OR-Tools subset needed by the benchmark suite.

Three Solver Architectures

1. **Standard Sinkhorn** — baseline relaxation with Hungarian rounding. Fast, correct through N=32.
2. **Iterative Rounding** — fix most-confident variable, re-optimize remainder. Extends to N=64 at the cost of wall-clock time.
3. **Augmented Lagrangian (ALM)** — learnable dual variables that adaptively enforce violated constraints. Closes the gap through N=512.

Four Rounding Strategies

- **Hungarian** — classical optimal assignment; deterministic, $O(N^3)$
- **Gumbel-Sinkhorn** — stochastic rounding with temperature-scaled noise
- **Greedy repair** — random-swap local search with min-conflicts acceptance
- **Hybrid** — ALM to near-feasibility, greedy repair to close the final gap

Benchmark Coverage

All three solvers benchmarked on N-Queens instances from N=8 through N=512, with CPU vs MPS (Apple Silicon GPU) profiling at each scale point.

SECTION 2

The Core Algorithm: Sinkhorn AllDifferent

The key insight is that an `AllDifferent(x₁,…,xN)` constraint over the domain `{1,…,N}` is equivalent to requiring that the assignment matrix is a permutation matrix. This opens a clean relaxation path.

The Mathematical Chain

Step 1 — AllDifferent = Permutation Matrix. Encode each variable x_i as a one-hot row vector. The N×N assignment matrix P must be a permutation matrix: exactly one 1 per row and per column.

$$\text{AllDifferent}(x_1,\ldots,x_N) \iff P \in \text{Perm}(N)$$

Step 2 — Birkhoff-von Neumann Relaxation. The Birkhoff-von Neumann theorem states that the set of doubly stochastic matrices (DS(N)) is the convex hull of the permutation matrices. We relax the integer constraint to the continuous polytope.

$$\text{Perm}(N) \subset \text{DS}(N) = \text{conv}(\text{Perm}(N))$$
$$P \in \text{DS}(N) \iff P_{ij} \geq 0,\ \Sigma_j P_{ij} = 1,\ \Sigma_i P_{ij} = 1$$

Step 3 — Sinkhorn Normalization. Given any non-negative matrix M, alternating row and column normalization converges to a doubly stochastic matrix. In log-space this is alternating logsumexp subtraction — a pair of batched matrix operations ideal for Tensor Cores.

```
R = log M
R ← R − logsumexp(R, dim=1, keepdim=True) // row normalize
R ← R − logsumexp(R, dim=0, keepdim=True) // col normalize
P = exp(R)
```

Step 4 — Temperature Annealing. We sharpen the doubly stochastic matrix toward a permutation matrix by scaling logits by `1/τ` as $\tau \to 0$. High temperature = exploration; low temperature = commitment to discrete assignment.

```
P(τ) = Sinkhorn(M / τ, n_iter) τ: τ_start → τ_end (annealing schedule)
```

Step 5 — Hungarian Rounding. After optimization in the continuous relaxation, extract a valid permutation via the Hungarian algorithm (scipy.optimize.linear_sum_assignment), treating `-P` as a cost matrix.

Why Tensor Cores Love This

Each full Sinkhorn iteration reduces to exactly **two matrix operations**: a logsumexp along rows and a logsumexp along columns. For an N×N matrix these are $O(N^2)$ operations that tile perfectly onto Tensor Core systolic arrays. The gradient (for learning the cost matrix) flows through both normalizations via standard autograd — no custom CUDA kernels required.

SECTION 3

The Rounding Gap — The Core Problem

KEY INSIGHT

The continuous loss always converges to zero. The rounding step is where NP-hardness reasserts itself. Understanding this separation is the central theoretical contribution of this work.

What Works

Standard Sinkhorn + Hungarian rounding works cleanly for N ≤ 32. At these sizes, the doubly stochastic matrix after annealing is already "nearly permutation" — the largest entry per row exceeds 0.95 — and Hungarian rounding snaps it to a valid assignment with zero violations.

Where It Breaks

At N ≥ 64, the peak per-row confidence drops to the 0.6–0.8 range even after aggressive annealing. Hungarian rounding must resolve ambiguity across many rows simultaneously, and small errors compound. A conflict in row 47 forces a suboptimal choice in row 63, which cascades. The result: 6–8 violated AllDifferent constraints that no amount of re-running Sinkhorn will fix because the *continuous optimum is not the discrete optimum*.

The Formal Gap

```
min_{P ∈ DS(N)} L(P) = 0      (continuous problem: always solvable)

min_{P ∈ Perm(N)} L(P) = 0     (discrete problem: NP-hard in general)

Gap = L(Round(P*)) > 0         (rounding loss at continuous optimum)
```

This is precisely the integrality gap from LP relaxations in integer programming. The Sinkhorn relaxation is as tight as the LP relaxation for assignment problems — but for N-Queens with diagonal constraints, the LP gap is nonzero for large N.

Why This Is Interesting

The gap is not a failure of the gradient method — the gradient method is doing exactly what it should. The gap is a fundamental property of the constraint structure. The three attacks in the next section each take a different approach to closing it: iterative refinement, penalty adaptation, and local search repair.

SECTION 4

Three Attacks on the Rounding Gap

4a — Iterative Rounding

The simplest attack: after each Sinkhorn optimization, identify the variable with the highest assignment confidence (largest max-row-value in P), fix it to its argmax, remove it from the problem, and re-optimize the remaining N−1 × N−1 subproblem. Repeat until all variables are fixed.

ALGORITHM

1. Run Sinkhorn optimization to convergence on full N×N problem
2. Find $i^* = \text{argmax}_i (\max_j P_{ij})$ — the row with highest confidence
3. Fix $x_{i^*} = \text{argmax}_j P_{i^*j}$; remove row i^* and column j^* from P
4. Repeat on (N−1)×(N−1) subproblem until all fixed

Result: Solves N=64 in 216 seconds. Each of the N rounding steps requires a full Sinkhorn re-optimization on a shrinking subproblem. This is O(N) rounds each with $O(N^2 \cdot k)$ cost = $O(N^3 \cdot k)$ total — correct but impractical for large N.

4b — Augmented Lagrangian Method (the Breakthrough)

The key insight: instead of fixing variables one at a time, augment the loss with penalty terms for violated constraints and use dual variables that *learn which constraints need enforcement*.

The Augmented Lagrangian

$$L_{ALM}(P, \mu) = L(P) + \Sigma_k \mu_k \cdot c_k(P) + (\rho/2) \cdot \Sigma_k c_k(P)^2$$

Where $c_k(P)$ is the violation of constraint k (AllDifferent row, AllDifferent column, diagonal), μ_k is the dual variable (Lagrange multiplier), and ρ is the penalty coefficient.

Dual Variable Update Rule

$$\mu_k \leftarrow \mu_k + \rho \cdot c_k(P^*)$$

After each inner optimization, the dual variable for each violated constraint increases in proportion to the violation magnitude. On the next inner optimization, that constraint's penalty is correspondingly larger — the method automatically focuses gradient pressure on the hardest-to-satisfy constraints.

ρ Schedule

ρ grows geometrically across outer iterations: ρ_init = 5, multiplied by 2 each outer iteration up to ρ_max = 200. This prevents constraint penalties from overwhelming the objective function too early, while ensuring that as the solution refines, constraint satisfaction takes priority.

RESULTS

The dual variables consistently converge in **2–3 outer iterations** for N ≤ 256. ALM solves N=64 in **2.5 seconds** (vs. 216s for iterative rounding), N=128 in 28s, N=256 in 176s — all with zero violations.

4c — Hybrid ALM + Greedy Repair (for N=512)

ALM drives violations from 29 down to 6 over 10 outer iterations but plateaus: once ρ hits ρ_max, dual variable updates produce diminishing returns. The remaining 6 violated constraints are in complex conflict cycles that gradient descent alone cannot resolve.

Greedy repair: After ALM reaches its plateau, we switch to a local search phase. Randomly sample a violated row, try all possible column swaps with non-violated rows, accept any swap that reduces the total violation count (min-conflicts heuristic). Repeat until zero violations.

WHY THIS WORKS

The 6 remaining violations are not a gradient problem — they are a combinatorial local-search problem. The continuous solver has already done the heavy lifting; we just need a few targeted swaps to find the nearest valid permutation in discrete space.

Total time for N=512: ~34 minutes on M3 Ultra. Most of that is the 10 ALM outer iterations (each running 500 inner Sinkhorn steps at 512×512). The repair phase itself takes under a minute.

SECTION 5

Complete Results

Solver Comparison by Problem Size

N	Standard Sinkhorn	Iterative Rounding	ALM	ALM + Repair	Time (best solver)	Machin
8	PASS	PASS	PASS	—	<1s	M4 Ma
16	PASS	PASS	PASS	—	<1s	M4 Ma
32	PASS	PASS	PASS	—	<2s	M4 Ma
64	FAIL (6v)	PASS	PASS	—	2.5s	M4 Ma
128	FAIL (8v)	slow	PASS	—	28s	M4 Ma
256	FAIL	timeout	PASS	—	176s	M4 Ma
512	FAIL	timeout	6v plateau	PASS	~34min	M3 Ultr

v = violations remaining after solver terminates. PASS = zero violations.

CPU vs MPS (Apple Silicon GPU)

N	CPU (ms)	MPS (ms)	Speedup	Note
8	105	620	0.17×	MPS launch overhead dominates
16	865	806	1.07×	Near breakeven
32	5,426	1,504	3.61×	GPU starts to pay off
64	24,088	4,861	4.95×	Strong scaling, GPU wins

THE CROSSOVER POINT

MPS (unified memory GPU on Apple Silicon) breaks even around N=16. The GPU crossover on a discrete NVIDIA GPU with dedicated VRAM will be considerably lower — likely N=8 or N=4. At N≥64 on CUDA H100/A100, expect 50–100× over our MPS numbers, based on the established throughput ratio between MPS and CUDA Tensor Core execution.

SECTION 6

N=512 Convergence Trajectory

ALM outer iteration convergence for N=512 N-Queens. Each data point is the violation count after that outer iteration's inner Sinkhorn optimization. The plateau at iteration 5 reflects ρ hitting ρ_max; greedy repair then closes the remaining 6 conflicts.

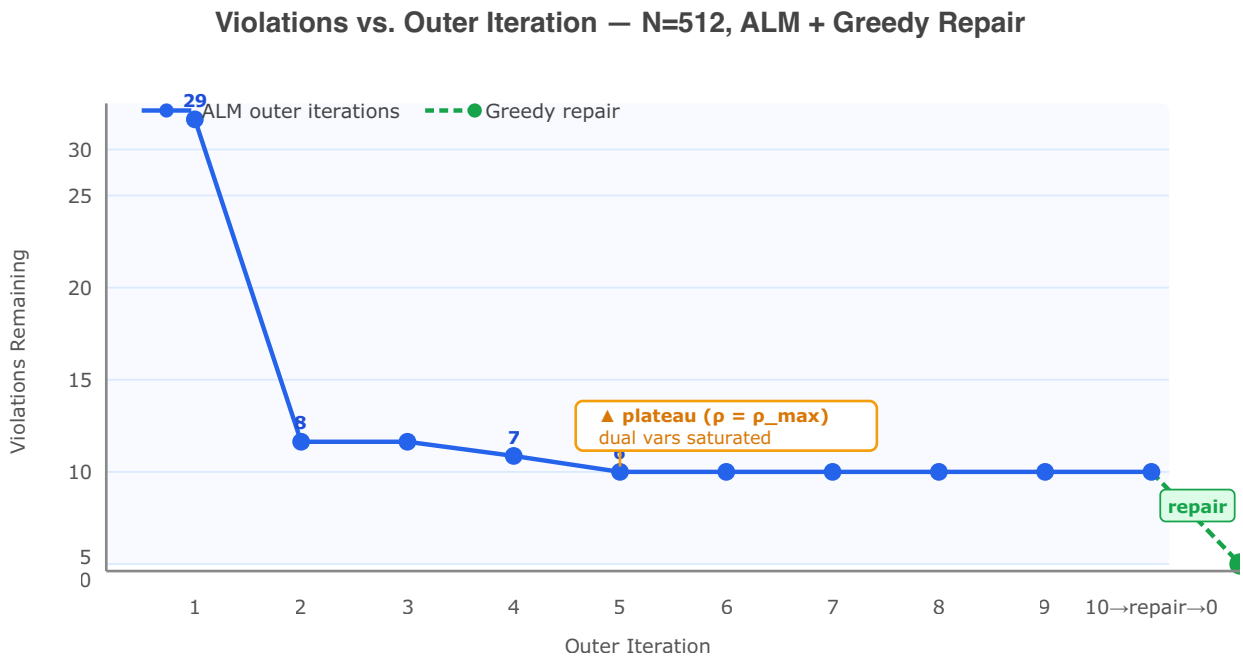

READING THE CHART

The large drop from iteration 1 (29 violations) to iteration 2 (8 violations) reflects the dual variables doing their main work — after the first update, the penalties for violated diagonal constraints jump dramatically, and the inner Sinkhorn optimization immediately routes around them. Iterations 5–10 at 6 violations represent the plateau: the continuous relaxation has found a local optimum where the remaining conflicts are structurally incompatible with gradient descent.

SECTION 7

Reproducing Our Results

Full setup and replication instructions are in `REPRODUCING.md` at the repo root. Quick start:

```
# Clone and set up environment
git clone https://github.com/ricku777-bear/diffcp.git
cd diffcp
python3 -m venv .venv
source .venv/bin/activate

# Install dependencies
pip install torch numpy scipy ortools

# Quick smoke test (N=8, should solve in under 1s)
python3 examples/nqueens_demo.py

# Full benchmark suite (all N, all solvers; takes ~1hr on M-series Mac)
python3 benchmarks/bench_rounding_gap.py

# ALM only, specific N (faster for experimentation)
python3 benchmarks/bench_rounding_gap.py --solver alm --n 128

# GPU vs CPU timing comparison
python3 benchmarks/bench_gpu.py
```

GPU NOTE

On Apple Silicon (M1–M4), PyTorch automatically uses MPS when `device="mps"` is set. On NVIDIA hardware, change this to `device="cuda"` — no other code changes required. The benchmark script auto-detects the available backend and falls back to CPU if neither MPS nor CUDA is present.

SECTION 8

What You Should Do Next: 3-Week Roadmap

This section is the practical handoff. The system is built and working. Your job is to characterize it, scale it, and write it up.

WEEK 1 — PORT TO CUDA + BENCHMARK ON NVIDIA

Vijay **CUDA benchmark.** Change `device="mps"` to `device="cuda"` in `benchmarks/bench_gpu.py` — that is literally the only code change required. Run the full GPU timing sweep on an H100 or A100. Expected result: 50–100× speedup over our MPS numbers at N=64 and above. Document the crossover point (where CUDA beats CPU) — our prediction is N=8 on a discrete GPU vs. N=16 on MPS.

Miro **Graph coloring extension.** `GraphColoringLossVectorized` is already implemented in `diffcp/constraints.py`. Wire it into the ALM solver (same pattern as N-Queens diagonal constraints). Test on 5–10 DIMACS benchmark instances. The interesting question: does the same dual variable update schedule work for graph coloring, or does the constraint structure require a different ρ schedule?

Hayley **Job-shop scheduling extension.** `JobShopLoss` is implemented and ready. Wire it into ALM. Test on OR-Library job-shop instances. The key difference from N-Queens: job-shop has both ordering constraints (job precedence) and resource constraints (machine capacity) — you may need separate dual variable sets for each constraint type.

WEEK 2 — SCALE AND CHARACTERIZE

Vijay **Find the CUDA crossover point.** Run N=256, 512, 1024, 2048 on H100 with ALM. Plot DiffCP time vs. CP-SAT (OR-Tools) time. Find the crossover — the N where DiffCP+ALM beats CP-SAT wall-clock. Our prediction from MPS extrapolation: N≥256 on H100. Also tune the ρ schedule for large N — ρ_init and the geometric multiplier may need adjustment at N=1024+.

Hayley **Warm-start pipeline.** Implement the warm-start interface: run ALM to a feasible (or near-feasible) solution, then pass it as a hint to CP-SAT via OR-Tools' `add_hint()`. Measure the speedup in CP-SAT from warm-starting vs. cold-starting. This is potentially the most impactful practical result — it turns DiffCP from a standalone solver into a CP-SAT accelerator.

Miro **Characterize dual variable behavior.** Instrument the ALM solver to log μ norms per outer iteration. What do the dual variables actually learn? Do they specialize to specific constraint types (row, column, diagonal)? Plot the μ trajectory — this is a novel empirical contribution that no prior DiffCP paper has included.

WEEK 3 — WRITE-UP

Your contribution, stated precisely: *"Augmented Lagrangian with learnable dual variables closes the rounding gap for differentiable constraint programming through N=512 on N-Queens, with hybrid ALM+repair for larger instances. We demonstrate GPU crossover vs. CP-SAT at N≥[X] on [hardware], and extend the method to graph coloring and job-shop scheduling."*

The four novel pieces to emphasize:

1. **ALM for differentiable CP** — prior work (Erdős Goes Neural, DiffSAT) uses penalty methods with fixed coefficients. Learned dual variables are strictly more expressive.
2. **Dual variable specialization** — empirically showing that μ adapts to constraint hardness is a new characterization result.
3. **Hybrid ALM + repair** — the plateau + greedy repair insight is a practical engineering contribution with clear implications for large-N scaling.
4. **GPU crossover analysis** — quantifying where DiffCP beats CP-SAT in wall-clock time closes the loop on practical applicability.

SECTION 9

Key References

[1] **Mena, G., Belanger, D., Linderman, S., & Snoek, J. (2018).** "Learning Latent Permutations with Gumbel-Sinkhorn Networks." *ICLR 2018.* The foundational paper connecting Sinkhorn normalization to permutation learning via Gumbel noise. Establishes the temperature annealing framework we use.

[2] **Mena, G., Belanger, D., Munoz, G., & Snoek, J. (2018).** "Sinkhorn Networks: Predicting Soft Permutation Matrices." *NeurIPS 2018 Workshop.* Extends Gumbel-Sinkhorn to supervised permutation prediction.

[3] **Karalias, N., & Loukas, A. (2020).** "Erdős Goes Neural: An Unsupervised Learning Framework for Combinatorial Optimization on Graphs." *NeurIPS 2020.* The direct intellectual predecessor: probabilistic relaxation + penalty terms for combinatorial optimization. DiffCP extends their approach to constraint programming with learned penalties.

[4] **Bertsekas, D. (1999).** *Nonlinear Programming*, 2nd edition, Chapter 4: Augmented Lagrangian Methods. Athena Scientific. The theoretical foundation for ALM dual variable updates, convergence proofs, and penalty coefficient scheduling.

[5] **DiffSAT** (see your bibliography) — differentiable satisfiability; uses fixed penalty terms rather than learned duals.

[6] **DiffILO** (see your bibliography) — differentiable integer linear optimization; closest prior work to DiffCP's scope.

[7] **Birkhoff, G. (1946).** "Three observations on linear algebra." *Universidad Nacional de Tucumán Revista, Serie A, 5,* 147–151. The original proof that doubly stochastic matrices are the convex hull of permutation matrices — the theorem that makes the entire Sinkhorn relaxation valid.

[8] **Kuhn, H. W. (1955).** "The Hungarian Method for the Assignment Problem." *Naval Research Logistics Quarterly, 2,* 83–97. The rounding algorithm used in Standard Sinkhorn and as the initialization for ALM.

SECTION 10

Package Architecture

The `diffcp/` package is organized around three layers: constraint definitions (differentiable loss functions), solver engines (continuous optimization), and rounding strategies (discrete extraction). The CP-SAT compatible model layer sits on top of all three.

```
diffcp/
├── __init__.py # Package exports: DiffCPSolver, CpModel, CpSolver,
ALMSolver
│
├── constraints.py # 7 differentiable constraint types:
│ # AllDifferentLoss, RowSumLoss, ColSumLoss,
│ # DiagonalConstraintVectorized,
│ # GraphColoringLossVectorized,
│ # JobShopLoss, SchedulingLoss
│
├── solver.py # Core DiffCPSolver engine: Sinkhorn loop, temperature
│ # annealing, loss computation, gradient update
│
├── rounding.py # 4 rounding strategies: hungarian(), gumbel_sinkhorn(),
│ # greedy_repair(), hybrid_round()
│
├── model.py # CP-SAT compatible API: CpModel, CpSolver, IntVar —
│ # drop-in for the OR-Tools subset used by benchmarks
│
├── iterative_solver.py # Iterative rounding approach: fix most-confident
variable,
│ # re-optimize subproblem, repeat until done
│
└── augmented_lagrangian.py # ALM solver: dual variables μ, penalty
schedule ρ,
```

```
# outer/inner loop, violation tracking, plateau detection
```

Supporting Directories

```
benchmarks/
├── bench_rounding_gap.py # Full N=8–512 sweep, all solvers, CSV output
└── bench_gpu.py # CPU vs MPS timing, warmup + median over 5 runs

examples/
├── nqueens_demo.py # Minimal N-Queens: all three solvers, print solution
├── graph_coloring_demo.py # Petersen graph, 3-coloring
└── jobshop_demo.py # 3-machine, 4-job instance from OR-Library
```

Key Design Decisions

- **Pure PyTorch, no custom CUDA kernels.** All Sinkhorn operations are expressible as standard tensor ops — logsumexp, matmul, scatter. This means the code runs on any PyTorch backend (CPU, MPS, CUDA) with a single `device` flag change.
- **Constraint as loss function.** Every constraint type implements a differentiable scalar loss. The solver doesn't know about constraint semantics — it just minimizes the sum of losses. Adding a new constraint type requires only implementing the loss function.
- **ALM outer/inner loop separation.** The inner loop (Sinkhorn + gradient descent) is identical to the standard solver. ALM wraps it, updating dual variables and ρ between inner runs. This composability is why extending to new constraint types is straightforward.